1850c9db-d2ec-4068-b04f-4fb7781005ef

<prev_user>

<request>

You have been given comments on the previous commits you made in the repository.

You are working on an issue in the 'ROCm/ATOM' repository.

<repository_context>
I've cloned the repository in the directory /home/runner/work/ATOM/ATOM (not in /tmp/inputs). Always use absolute paths when referring to files in the repository.
</repository_context>

<current_datetime>2026-01-23T05:47:44.277Z</current_datetime>

Consider the following problem statement:

<problem_statement>

This section includes details on the pull request with code changes you have already submitted to fix the problem above.

<pr_title>
CI: Add pre-check status check and fix code style issues
</pr_title>

<pr_description>

</pr_description>

Comments on the PR (you are @copilot in this section)

<comments>

<pr_comments>

</pr_comments>

<pr_reviews>

<comment_thread_resolved>
<file>.github/workflows/pre-checks.yaml:19</file>
<original_commit>85ffd48</original_commit>

</comment_thread_resolved>

<comment_thread_resolved>
<file>.github/workflows/pre-checks.yaml:29</file>
<original_commit>85ffd48</original_commit>

</comment_thread_resolved>

<comment_thread_resolved>
<file>.github/workflows/atom-test.yaml:22</file>
<original_commit>85ffd48</original_commit>

</comment_thread_resolved>

</pr_reviews>

<pr_comments>

</pr_comments>

<pr_reviews>

<comment_thread_resolved>
<file>atom/benchmarks/benchmark_serving.py:808</file>
<original_commit>d4cdfe9</original_commit>

</comment_thread_resolved>

</pr_reviews>

<pr_comments>

</pr_comments>

<pr_reviews>

<comment_thread_resolved>
<file>atom/benchmarks/benchmark_serving.py:808</file>
<original_commit>14c8d98</original_commit>

</comment_thread_resolved>

</pr_reviews>

<pr_comments>

</pr_comments>

<pr_reviews>

<comment_thread_resolved>
<file>.github/workflows/atom-test.yaml:258</file>
<original_commit>a27908a</original_commit>

```
  </comment_thread_resolved>

  </pr_reviews>

  <pr_comments>



  </pr_comments>

  <pr_reviews>

  <comment_thread_resolved>
  <file>.github/workflows/pre-checks.yaml:18-20</file>
  <original_commit>8e5a172</original_commit>



  </comment_thread_resolved>

  <comment_thread_resolved>
  <file>.github/workflows/atom-test.yaml:114</file>
  <original_commit>8e5a172</original_commit>



  </comment_thread_resolved>

  <comment_thread_resolved>
  <file>.github/scripts/atom_test.sh:17</file>
  <orig/pr_comments>

  </comments>

  ----
  The last **33** git commits in this branch are the changes you
  have made so far. Use those as your change commit history.
  </problem_statement>


Analyze and determine if action or explanation is needed
<rules_for_making_changes>
* Start by fully understanding <problem_statement> and the
<comments> before making any changes. You will only address the
<comment_new> comments. Reply to new comments as you address them.
```

Ignore comments that are not actionable or directed at you.
* For each new comment, determine:
    - If it a request for changes, a question on the code or a suggestion.
    - Is it addressed to you or someone else, e.g. "Hey @username, ..." or "@username can you ...".
    - Is it a general comment or praise.
* If there are no comments you need to address or reply to, you can stop now without exploring or explaining yourself.
* If the user's ask is not clear, you can reply to the comment and ask for clarification.
* If an action or explanation is needed, follow the steps below.
</rules_for_making_changes>

A **CI/Build Failure** has been mentioned in the problem statement, make sure to follow your instructions for addressing CI/Build Failures.

<prompt>

@copilot Please review my code and analysis the cause of this error:
```
Run # Run the inference and capture output
========== Running test ==========
Model path: amd/Llama-3.3-70B-Instruct-MXFP4-Preview
ls: cannot access 'amd/Llama-3.3-70B-Instruct-MXFP4-Preview': No such file or directory
[aiter] start build [module_aiter_enum] under /app/aiter-test/aiter/jit/build/module_aiter_enum
[aiter] Current hipcc not support: -mllvm -amdgpu-coerce-illegal-types=1, skip it.
[aiter] finish build [module_aiter_enum], cost 7.1s
[aiter] import [module_aiter_enum] under /app/aiter-test/aiter/jit/module_aiter_enum.so
`torch_dtype` is deprecated! Use `dtype` instead!
[atom] Engine Core Mgr: Creating EngineCore for DP rank 0/1
[atom] Creating EngineCore process: DP rank 0, will use GPUs 0 to 0
[atom] Engine Core Mgr: Starting EngineCore for DP rank 0/1
[aiter] message queue communication handle: Handle(local_reader_ranks=[0], buffer_handle=(1, 16777216, 10, 'psm_3376a6f1'), local_subscribe_addr='ipc:///tmp/
```

```
6885b872-5a12-40fd-b9c5-a58d88d2c6d2', remote_subscribe_addr=None,
remote_addr_ipv6=False)
[aiter] import [module_aiter_enum] under /app/aiter-test/aiter/jit/
module_aiter_enum.so
[aiter] merge tuned file under model_configs/ and configs/ /app/
aiter-test/aiter/configs/bf16_tuned_gemm.csv:/app/aiter-test/aiter/
configs/model_configs/gptoss_bf16_tuned_gemm.csv:/app/aiter-test/
aiter/configs/model_configs/dsv3_bf16_tuned_gemm.csv
`torch_dtype` is deprecated! Use `dtype` instead!
[atom] ModelRunner rank=0, dp_rank_local=0, local_device_rank=0,
device=cuda:0
[aiter] rank 0 in world size 1 is assigned as DP rank 0, PP rank 0,
TP rank 0, EP rank 0
[atom.model_loader.weight_utils] Using model weights format
['*.safetensors']
[atom.model_loader.weight_utils] Time spent downloading weights for
amd/Llama-3.3-70B-Instruct-MXFP4-Preview: 67.289874 seconds
Loading safetensors shards[amd/Llama-3.3-70B-Instruct-MXFP4-
Preview]:   0%|          | 0/9 [00:00<?, ?it/s]
Loading safetensors shards[amd/Llama-3.3-70B-Instruct-MXFP4-
Preview]:  11%|█         | 1/9 [00:00<00:01,  7.97it/s]
Loading safetensors shards[amd/Llama-3.3-70B-Instruct-MXFP4-
Preview]: 100%|██████████| 9/9 [00:00<00:00, 51.18it/s]
[atom] Model load done: amd/Llama-3.3-70B-Instruct-MXFP4-Preview
[atom] Model Runner0/1: DP size=1 too large, warmup_max_tokens=16384
< max_model_len=131072. Using 1 seq with length 16384 for warmup.
[atom] AsyncIOProcManager(ModelRunner): [ModelRunner0/1] proc died
unexpectedly, shutting down.
[atom] AsyncIOProcManager(ModelRunner): shutdown all runners...
[atom] AsyncIOProcManager(ModelRunner): All runners are shutdown.
[atom] Engine Core: load model runner failed
[atom] Engine Core Mgr: Failed to initialize all EngineCores,
cleaning up: Engine Core Mgr: Expected READY signal from DP rank 0,
but got EngineCoreRequestType.ADD
[atom] Engine Core Mgr: Shutting down all 1 EngineCores
[atom] Engine Core Mgr: All EngineCores shut down
Traceback (most recent call last):
  File "<frozen runpy>", line 198, in _run_module_as_main
  File "<frozen runpy>", line 88, in _run_code
  File "/workspace/atom/examples/simple_inference.py", line 83, in
<module>
    main()
  File "/workspace/atom/examples/simple_inference.py", line 48, in
main
    llm = engine_args.create_engine()
          ^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "/workspace/atom/model_engine/arg_utils.py", line 206, in
```

```
create_engine
    return LLMEngine(self.model, **self._get_engine_kwargs())
           ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "/workspace/atom/model_engine/llm_engine.py", line 41, in
__init__
    self.core_mgr = CoreManager(config)
                    ^^^^^^^^^^^^^^^^^^^
  File "/workspace/atom/model_engine/engine_core_mgr.py", line 121,
in __init__
    self._wait_for_all_ready_signals()
  File "/workspace/atom/model_engine/engine_core_mgr.py", line 188,
in _wait_for_all_ready_signals
    raise RuntimeError(
RuntimeError: Engine Core Mgr: Expected READY signal from DP rank 0,
but got EngineCoreRequestType.ADD
/usr/lib/python3.12/multiprocessing/resource_tracker.py:254:
UserWarning: resource_tracker: There appear to be 1 leaked semaphore
objects to clean up at shutdown
  warnings.warn('resource_tracker: There appear to be %d '
/usr/lib/python3.12/multiprocessing/resource_tracker.py:254:
```

<prev_assistant>

Summary

I've analyzed the CI test failure for the Llama-3.3-70B-Instruct-MXFP4-Preview model and identified two issues:

1. **Fixed**: The `ls` command was executing on the host runner instead of inside the Docker container, causing it to fail when the `/models` directory wasn't available. I moved the `ls` command inside the Docker container (commit c90f3a7).

2. **Identified but not fixed**: The actual test failure is a model runner crash during warmup that appears to be an ATOM runtime issue, not related to the workflow changes. The logs show the model loads successfully but crashes during the warmup phase. This requires investigation in the ATOM codebase itself, which is outside the scope of the workflow fixes.

The workflow fix improves debugging output by running the `ls` command in the correct context where the model will actually be accessed.

<current>

<request>

You have been given comments on the previous commits you made in the repository.

You are working on an issue in the 'ROCm/ATOM' repository.

<repository_context>
I've cloned the repository in the directory /home/runner/work/ATOM/ATOM (not in /tmp/inputs). Always use absolute paths when referring to files in the repository.
</repository_context>

<current_datetime>2026-01-23T05:54:43.167Z</current_datetime>

Consider the following problem statement:

<problem_statement>

This section includes details on the pull request with code changes you have already submitted to fix the problem above.

<pr_title>
Fix ls command execution context in CI workflow
</pr_title>

<pr_description>
The CI workflow was executing `ls -la $model_path` on the host runner before running the Docker container inference test. When the `/models` directory doesn't exist, `model_path` is set to a HuggingFace model identifier (e.g., `amd/Llama-3.3-70B-Instruct-MXFP4-Preview`), which only exists inside the container after download, causing the `ls` command to fail.

Changes

- Move `ls` command inside the Docker container execution block where the model path is actually accessible
- Add informative fallback message when model doesn't exist yet (before HuggingFace download)

```yaml
# Before: runs on host
echo "Model path: $model_path"
```

```
ls -la $model_path || true
docker exec atom_test bash -lc "python3 -m
atom.examples.simple_inference..."

# After: runs inside container
echo "Model path: $model_path"
docker exec atom_test bash -lc "
  ls -la \"$model_path\" || echo \"Model path does not exist yet
(will be downloaded by HuggingFace)\"
  python3 -m atom.examples.simple_inference...
"
```

This ensures the debug output runs in the same context where the model will be accessed, providing accurate diagnostics regardless of whether `/models` is pre-mounted or models are downloaded on-demand.
</pr_description>

Comments on the PR (you are @copilot in this section)

<comments>

<pr_comments>

<comment_new>
<comment_id>3695812742</comment_id>
<author>@gyohuangxin</author>
What about this error?
```
[atom] Model load done: amd/Llama-3.3-70B-Instruct-MXFP4-
Preview
[atom] Model Runner0/1: DP size=1 too large,
warmup_max_tokens=16384 < max_model_len=131072. Using 1 seq
with length 16384 for warmup.
[atom] AsyncIOProcManager(ModelRunner): [ModelRunner0/1]
proc died unexpectedly, shutting down.
[atom] AsyncIOProcManager(ModelRunner): shutdown all
runners...
[atom] AsyncIOProcManager(ModelRunner): All runners are
shutdown.
[atom] Engine Core: load model runner failed
[atom] Engine Core Mgr: Failed to initialize all
EngineCores, cleaning up: Engine Core Mgr: Expected READY
signal from DP rank 0, but got EngineCoreRequestType.ADD
[atom] Engine Core Mgr: Shutting down all 1 EngineCores
[atom] Engine Core Mgr: All EngineCores shut down
Traceback (most recent call last):
```

```
  File "<frozen runpy>", line 198, in _run_module_as_main
  File "<frozen runpy>", line 88, in _run_code
  File "/workspace/atom/examples/simple_inference.py", line
83, in <module>
    main()
  File "/workspace/atom/examples/simple_inference.py", line
48, in main
    llm = engine_args.create_engine()
          ^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "/workspace/atom/model_engine/arg_utils.py", line
206, in create_engine
    return LLMEngine(self.model,
**self._get_engine_kwargs())

^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
  File "/workspace/atom/model_engine/llm_engine.py", line
41, in __init__
    self.core_mgr = CoreManager(config)
                    ^^^^^^^^^^^^^^^^^^^
  File "/workspace/atom/model_engine/engine_core_mgr.py",
line 121, in __init__
    self._wait_for_all_ready_signals()
  File "/workspace/atom/model_engine/engine_core_mgr.py",
line 188, in _wait_for_all_ready_signals
    raise RuntimeError(
RuntimeError: Engine Core Mgr: Expected READY signal from DP
rank 0, but got EngineCoreRequestType.ADD
/usr/lib/python3.12/multiprocessing/resource_tracker.py:254:
UserWarning: resource_tracker: There appear to be 1 leaked
semaphore objects to clean up at shutdown
  warnings.warn('resource_tracker: There appear to be %d '
/usr/lib/python3.12/multiprocessing/resource_tracker.py:254:
```
</comment_new>

</pr_comments>

</comments>

The last **2** git commits in this branch are the changes you have made so far. Use those as your change commit history.
</problem_statement>

Analyze and determine if action or explanation is needed
<rules_for_making_changes>
* Start by fully understanding <problem_statement> and the

<comments> before making any changes. You will only address the <comment_new> comments. Reply to new comments as you address them. Ignore comments that are not actionable or directed at you.
* For each new comment, determine:
    - If it a request for changes, a question on the code or a suggestion.
    - Is it addressed to you or someone else, e.g. "Hey @username, ..." or "@username can you ...".
    - Is it a general comment or praise.
* If there are no comments you need to address or reply to, you can stop now without exploring or explaining yourself.
* If the user's ask is not clear, you can reply to the comment and ask for clarification.
* If an action or explanation is needed, follow the steps below.
</rules_for_making_changes>

A **CI/Build Failure** has been mentioned in the problem statement, make sure to follow your instructions for addressing CI/Build Failures.