fed94b0c-2e43-46b4-8804-c2dd6c94808e

<prev_user>

<request>

You have been given comments on the previous commits you made in the repository.

You are working on an issue in the 'erigontech/erigon' repository.

<repository_context>
I've cloned the repository in the directory /home/runner/work/erigon/erigon (not in /tmp/inputs). Always use absolute paths when referring to files in the repository.
</repository_context>

<current_datetime>2026-05-07T11:04:38.829+00:00</current_datetime>

Consider the following problem statement:

<problem_statement>

This section includes details on the pull request with code changes you have already submitted to fix the problem above.

<pr_title>
ci(docs): auto-update hardware requirements disk sizes from sync CI
</pr_title>

<pr_description>
Summary

Introduces an automated pipeline to keep the disk size figures on the Hardware Requirements page up to date, removing the need for manual edits.

How it works

1. **Measure** — `qa-sync-from-scratch` and `qa-sync-from-scratch (minimal node)` each run a `du -sb` step just before their cleanup, uploading a `disk-usage<chain><mode>.txt` artifact.
2. **Collect** — A new `update-disk-sizes.yml` workflow triggers via `workflow_run` on successful completion of either sync workflow, downloads all matching artifacts, and runs `docs/site/scripts/update-disk-sizes.py` to merge the bytes into `docs/site/src/data/disk-sizes.json`.
3. **Publish** — The workflow commits the updated JSON to `docs/auto/disk-sizes` and opens (or updates) a **draft PR** against `release/3.4` for human review before merging.
4. **Render** — `hardware-requirements.mdx` now imports `disk-sizes.json` and renders the *Current Disk Usage* column dynamically — no manual edits to the MDX needed after the first CI run.

What is automated

Network	Mode	Source
Ethereum mainnet	Full	`qa-sync-from-scratch` (weekly)
Ethereum mainnet	Minimal	`qa-sync-from-scratch (minimal node)` (nightly)
Gnosis Chain	Full	`qa-sync-from-scratch` (weekly)
Gnosis Chain	Minimal	`qa-sync-from-scratch (minimal node)` (nightly)

What is NOT automated yet — follow-up

Archive mode disk sizes require measurement from always-on snapshot machines rather than ephemeral CI runners (a full archive sync from scratch is impractical as a CI step). A follow-up PR will add this once snapshot machine access/outbound push is confirmed with the DevOps team.

The JSON also accepts `workflow_dispatch` with a custom `run_id` + `prune_mode`, so measurements can be triggered manually from any run.

Files changed

File	Change
`.github/workflows/qa-sync-from-scratch.yml`	Add measure + upload steps before cleanup
`.github/workflows/qa-sync-from-scratch-minimal-node.yml`

Same |
`.github/workflows/update-disk-sizes.yml`	New collector workflow
`docs/site/src/data/disk-sizes.json`	New data file (seeded with current Sept 2025 values)
`docs/site/scripts/update-disk-sizes.py`	New Python helper — parses artifact bytes, formats SI units, updates JSON
`docs/site/docs/get-started/hardware-requirements.mdx`	Import JSON, render disk usage cells dynamically, remove static dated note

🤖 Generated with Claude Code
</pr_description>

<stacked_pr_context>
The pull request you are currently working on is a stacked pull request on top of another pull request. The base branch is `release/3.4`, not the default branch `main`. A pull request has already been created for your branch. Do not attempt to create a new pull request. If you are asked to fix CI, build, test, or workflow failures, check the GitHub Actions workflow runs for the original branch (`release/3.4`) first. If you don't find relevant failures there, then check your branch (`docs/auto-disk-sizes`).
</stacked_pr_context>

Comments on the PR (you are @copilot in this section)

<comments>

<pr_comments>

</pr_comments>

<pr_reviews>

<comment_thread>
<file>.github/workflows/update-disk-sizes.yml:45-49</file>
<original_commit>acd26a7</original_commit>

>

</comment_thread>

<comment_thread>
<file>.github/workflows/update-disk-sizes.yml:104-112</file>
<original_commit>acd26a7</original_commit>

<comment_old>
<author>@copilot-pull-request-reviewer</author>
The update branch is always recreated from the currently checked-out branch (`git checkout -B $BRANCH` after checking out `main`), even if `origin/$BRANCH` already exists. This means consecutive runs (e.g., minimal then full) will overwrite each other’s previous JSON updates, so the draft PR will only contain the last run’s prune-mode changes. To preserve/merge updates, base the working branch on `origin/$BRANCH` when it exists (or check out the branch before running the updater), and consider adding workflow concurrency to avoid simultaneous force-push races.
</comment_old>

<comment_new>
<comment_id>3200776219</comment_id>
<author>@bloxster</author>
@copilot apply changes based on this feedback
</comment_new>

</comment_thread>

<comment_thread_resolved>
<file>docs/site/docs/get-started/hardware-requirements.mdx:28-36</file>
<original_commit>acd26a7</original_commit>

<comment_old>
<author>@copilot-pull-request-reviewer</author>
After switching to JSON-driven values, the page no longer shows when “Current Disk Usage” was measured, and it still labels the (manual, non-auto-updated) archive numbers as “Current”. This can be misleading if CI hasn’t run recently or for archive entries that are explicitly not automated. Consider rendering `diskSizes.ci_last_updated` (and/or per-row `measured_at`/`source`) near the table, or add an explicit note that archive values are manual and may be stale.
</comment_old>

<comment_old>
<author>@bloxster</author>

The routine runs weekly, no need to worry.
</comment_old>

</comment_thread_resolved>

</pr_reviews>

</comments>

The last **3** git commits in this branch are the changes you have made so far. Use those as your change commit history.
</problem_statement><rules_for_making_changes>
* Start by fully understanding <problem_statement> and the <comments> before making any changes. You will only address the <comment_new> comments. Reply to new comments as you address them. Ignore comments that are not actionable or directed at you.
* For each new comment, determine:
    - If it a request for changes, a question on the code or a suggestion.
    - Is it addressed to you or someone else, e.g. "Hey @username, ..." or "@username can you ...".
    - Is it a general comment or praise.
* If there are no comments you need to address or reply to, you can stop now without exploring or explaining yourself.
* If the user's ask is not clear, you can reply to the comment and ask for clarification.
* If an action or explanation is needed, follow the steps below.
</rules_for_making_changes>

<system_reminder>
You have the capability to call multiple tools in a single response.
For maximum efficiency, whenever you need to perform multiple independent operations, ALWAYS call tools simultaneously whenever the actions can be done in parallel rather than sequentially (e.g. multiple reads/edits to different files). Especially when exploring repository, searching, reading files, viewing directories, validating changes. For example, you can read 3 different files in parallel, or edit different files in parallel. However, if some tool calls depend on previous calls to inform dependent values like the parameters, do NOT call

these tools in parallel and instead call them sequentially (e.g. reading shell output from a previous command should be sequential as it requires the sessionID).
</system_reminder>

A **CI/Build Failure** has been mentioned in the problem statement, make sure to follow your instructions for addressing CI/Build Failures.

<prev_assistant>

Done — I addressed the requested review feedback in `.github/workflows/update-disk-sizes.yml` and pushed the fixes.

What I changed
- Preserved prior updates by preparing `docs/auto/disk-sizes` from `origin/docs/auto/disk-sizes` when it already exists (instead of always recreating from base).
- Added workflow `concurrency` to reduce force-push race conditions between overlapping runs.
- Switched checkout/PR base to the source run branch (`github.event.workflow_run.head_branch`), with `release/3.4` as the manual dispatch fallback.
- Kept fetch error visibility while making branch-existence checks explicit.
- Fixed workflow formatting/indentation issues introduced during edits.

Validation performed
- Investigated Actions runs/logs via GitHub MCP on `release/3.4` first (then `docs/auto-disk-sizes`).
- Ran `make lint` before and after changes (clean).
- YAML parse validation passed.
- Ran `parallel_validation` (Code Review + CodeQL) after changes.

Comment response
- Replied to `comment_id: 3200776219` with commit reference: `2e81558`.

Remaining non-blocking review note
- One maintainability suggestion remains from automated review (deduplicate hardcoded branch name into a job-level env var). I left it as-is to keep this fix minimal and focused.

<current>

<request>

You have been given comments on the previous commits you made in the repository.

You are working on an issue in the 'erigontech/erigon' repository.

<repository_context>
I've cloned the repository in the directory /home/runner/work/erigon/erigon (not in /tmp/inputs). Always use absolute paths when referring to files in the repository.
</repository_context>

<current_datetime>2026-05-07T11:21:07.487+00:00</current_datetime>

Consider the following problem statement:

<problem_statement>

This section includes details on the pull request with code changes you have already submitted to fix the problem above.

<pr_title>
ci(docs): auto-update hardware requirements disk sizes from sync CI
</pr_title>

<pr_description>
Summary

Introduces an automated pipeline to keep the disk size figures on the Hardware Requirements page up to date, removing the need for manual edits.

How it works

1. **Measure** — `qa-sync-from-scratch` and `qa-sync-from-scratch (minimal node)` each run a `du -sb` step just before their cleanup, uploading a `disk-usage<chain><mode>.txt` artifact.
2. **Collect** — A new `update-disk-sizes.yml` workflow triggers via `workflow_run` on successful completion of either sync workflow, downloads all matching artifacts, and runs `docs/site/

```
scripts/update-disk-sizes.py` to merge the bytes into `docs/
site/src/data/disk-sizes.json`.
3. **Publish** — The workflow commits the updated JSON to `docs/
auto/disk-sizes` and opens (or updates) a **draft PR** against
`release/3.4` for human review before merging.
4. **Render** — `hardware-requirements.mdx` now imports `disk-
sizes.json` and renders the *Current Disk Usage* column
dynamically — no manual edits to the MDX needed after the first
CI run.

### What is automated

| Network | Mode | Source |
|---------|------|--------|
| Ethereum mainnet | Full | `qa-sync-from-scratch` (weekly) |
| Ethereum mainnet | Minimal | `qa-sync-from-scratch (minimal
node)` (nightly) |
| Gnosis Chain | Full | `qa-sync-from-scratch` (weekly) |
| Gnosis Chain | Minimal | `qa-sync-from-scratch (minimal node)`
(nightly) |

### What is NOT automated yet — follow-up

**Archive mode** disk sizes require measurement from always-on
snapshot machines rather than ephemeral CI runners (a full
archive sync from scratch is impractical as a CI step). A
follow-up PR will add this once snapshot machine access/outbound
push is confirmed with the DevOps team.

The JSON also accepts `workflow_dispatch` with a custom `run_id`
+ `prune_mode`, so measurements can be triggered manually from
any run.

### Files changed

| File | Change |
|------|--------|
| `.github/workflows/qa-sync-from-scratch.yml` | Add measure +
upload steps before cleanup |
| `.github/workflows/qa-sync-from-scratch-minimal-node.yml` |
Same |
| `.github/workflows/update-disk-sizes.yml` | New collector
workflow |
| `docs/site/src/data/disk-sizes.json` | New data file (seeded
with current Sept 2025 values) |
| `docs/site/scripts/update-disk-sizes.py` | New Python helper —
parses artifact bytes, formats SI units, updates JSON |
```

```
| `docs/site/docs/get-started/hardware-requirements.mdx` |
Import JSON, render disk usage cells dynamically, remove static
dated note |

🤖 Generated with <a href="https://claude.com/claude-
code">Claude Code</a>
</pr_description>


<stacked_pr_context>
The pull request you are currently working on is a stacked pull
request on top of another pull request. The base branch is
`release/3.4`, not the default branch `main`. A pull request has
already been created for your branch. Do not attempt to create a
new pull request. If you are asked to fix CI, build, test, or
workflow failures, check the GitHub Actions workflow runs for
the original branch (`release/3.4`) first. If you don't find
relevant failures there, then check your branch (`docs/auto-
disk-sizes`).
</stacked_pr_context>

## Comments on the PR (you are @copilot in this section)

<comments>

<pr_comments>

<comment_old>
<author>@copilot-pull-request-reviewer</author>
## Pull request overview

This PR introduces an automated CI-to-docs pipeline that
measures Erigon sync disk usage in QA workflows, writes the
results into a JSON data file, and renders the “Current Disk
Usage” column on the hardware requirements docs page from that
data.

**Changes:**
- Add disk-usage measurement + artifact upload steps to the full
and minimal “sync from scratch” QA workflows.
- Add a new `update-disk-sizes` workflow plus a Python script to
collect artifacts and update `docs/site/src/data/disk-
sizes.json`.
- Update the hardware requirements MDX to render disk usage from
`disk-sizes.json` instead of hard-coded values.

### Reviewed changes
```

Copilot reviewed 6 out of 6 changed files in this pull request and generated 3 comments.

💡 <a>Add Copilot custom instructions for smarter, more guided reviews. Learn how to get started.
</comment_old>

</pr_comments>

<pr_reviews>

<comment_thread>
<file>.github/workflows/update-disk-sizes.yml:45-49</file>
<original_commit>acd26a7</original_commit>

<comment_old>
<author>@copilot-pull-request-reviewer</author>
This workflow always checks out `main` and creates the update PR against `main` (`ref: main`, `--base main`). That conflicts with the PR description (targets `release/3.4`) and with the docs deployment workflow, which appears to publish from `release/3.4` (so updates merged to `main` may never reach the live docs). Consider checking out / targeting the same docs release branch (e.g. `release/3.4`, or `github.event.workflow_run.head_branch`) and using that as the PR base.
</comment_old>

<comment_new>
<comment_id>3200776181</comment_id>
<author>@bloxster</author>
@copilot apply changes based on this feedback
</comment_new>

```
</comment_thread>

<comment_thread>
<file>.github/workflows/update-disk-sizes.yml:121-127</file>
<original_commit>acd26a7</original_commit>

<comment_old>
<author>@copilot-pull-request-reviewer</author>
The update branch is always recreated from the currently
checked-out branch (`git checkout -B $BRANCH` after checking out
`main`), even if `origin/$BRANCH` already exists. This means
consecutive runs (e.g., minimal then full) will overwrite each
other’s previous JSON updates, so the draft PR will only contain
the last run’s prune-mode changes. To preserve/merge updates,
base the working branch on `origin/$BRANCH` when it exists (or
check out the branch before running the updater), and consider
adding workflow concurrency to avoid simultaneous force-push
races.
</comment_old>

<comment_old>
<author>@bloxster</author>
@copilot apply changes based on <a href="https://github.com/
erigontech/erigon/pull/21030#discussion_r3200776219">this
feedback</a>
</comment_old>

<comment_old>
<author>@copilot</author>
Implemented in 2e81558. The workflow now prepares `docs/auto/
disk-sizes` from `origin/docs/auto/disk-sizes` when it exists
(so updates accumulate instead of being overwritten), adds
workflow concurrency to avoid force-push races, and uses the
source run branch as checkout/PR base (with `release/3.4` for
manual dispatch).
</comment_old>

</comment_thread>

<comment_thread_resolved>
<file>docs/site/docs/get-started/hardware-
requirements.mdx:28-36</file>
<original_commit>acd26a7</original_commit>

<comment_old>
<author>@copilot-pull-request-reviewer</author>
After switching to JSON-driven values, the page no longer shows
```

when “Current Disk Usage” was measured, and it still labels the (manual, non-auto-updated) archive numbers as “Current”. This can be misleading if CI hasn’t run recently or for archive entries that are explicitly not automated. Consider rendering `diskSizes.ci_last_updated` (and/or per-row `measured_at`/`source`) near the table, or add an explicit note that archive values are manual and may be stale.
</comment_old>

<comment_old>
<author>@bloxster</author>
The routine runs weekly, no need to worry.
</comment_old>

</comment_thread_resolved>

</pr_reviews>

</comments>

The last **6** git commits in this branch are the changes you have made so far. Use those as your change commit history.
</problem_statement><rules_for_making_changes>
* Start by fully understanding <problem_statement> and the <comments> before making any changes. You will only address the <comment_new> comments. Reply to new comments as you address them. Ignore comments that are not actionable or directed at you.
* For each new comment, determine:
    - If it a request for changes, a question on the code or a suggestion.
    - Is it addressed to you or someone else, e.g. "Hey @username, ..." or "@username can you ...".
    - Is it a general comment or praise.
* If there are no comments you need to address or reply to, you can stop now without exploring or explaining yourself.
* If the user's ask is not clear, you can reply to the comment and ask for clarification.
* If an action or explanation is needed, follow the steps below.
</rules_for_making_changes>

<system_reminder>

You have the capability to call multiple tools in a single response.
For maximum efficiency, whenever you need to perform multiple independent operations, ALWAYS call tools simultaneously whenever the actions can be done in parallel rather than sequentially (e.g. multiple reads/edits to different files). Especially when exploring repository, searching, reading files, viewing directories, validating changes. For example, you can read 3 different files in parallel, or edit different files in parallel. However, if some tool calls depend on previous calls to inform dependent values like the parameters, do NOT call these tools in parallel and instead call them sequentially (e.g. reading shell output from a previous command should be sequential as it requires the sessionID).
</system_reminder>

A **CI/Build Failure** has been mentioned in the problem statement, make sure to follow your instructions for addressing CI/Build Failures.